

7. RECOMMENDER SYSTEMS

Application Areas

The screenshot displays three distinct application areas for recommender systems:

- Products:** Shows recommendations for items similar to the one being viewed. It includes three products: a Pandigital digital picture frame, a ViewSonic multimedia device, and a Foscam IP camera.
- People:** Shows suggestions for users the user may know. It lists several individuals with their profiles and connection counts.
- News:** Shows recommended news articles based on the user's reading history. Headlines include "Indonesia's GDP Misses Estimates, Growing Least Since 2009", "China Manufacturing Gauge Signals Risk of Deeper Slowdown", and "How Russia Inc. Moves Billions Offshore -- and a Handful of Tax Havens May Hold Key to Sanctions".

At the bottom left, there is a copyright notice: ©2018, Karl Aberer, EPFL-IC, Laboratoire de systèmes d'informations répartis. At the bottom right, there is a footer note: Applied Classification - 2.

Recommender systems are widely used by many websites and applications to provide value for the user and the provider.

For users recommender systems help, e.g. in the context of a product search,

- find things that are interesting
- narrow down the set of choices
- help explore the space of options
- Discover new things

For providers recommender systems help to

- increase trust and customer loyalty with personalized services
- increase sales, click rates, page views etc.
- identify opportunities for promotion and persuasion
- obtain more knowledge about customers

Definition: Recommender System

Given a *user* model and a set of *items*, a recommender system is a function that helps to match users with items by *ranking* the items in order of decreased *relevance*



©2018, Karl Aberer, EPFL-IC, Laboratoire de systèmes d'informations répartis

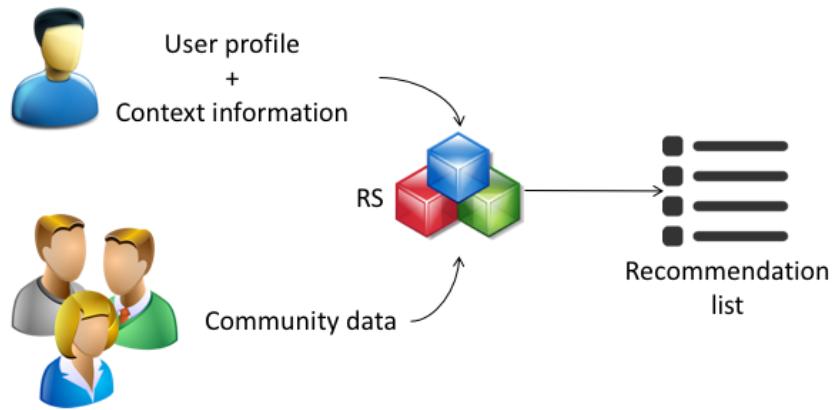
Applied Classification - 3

A recommender system identifies for a given user a set of items that are relevant to her. This can result in a ranking of the items (modelling a relevance function) or a classification.

The user model typically includes data such as ratings, preferences, demographics and attributes that characterize the user. The items can also be enriched with attributes about their content and characteristics, although it is not always necessary.

Collaborative Recommender System

Collaborative = “Tell me what **other people** like”



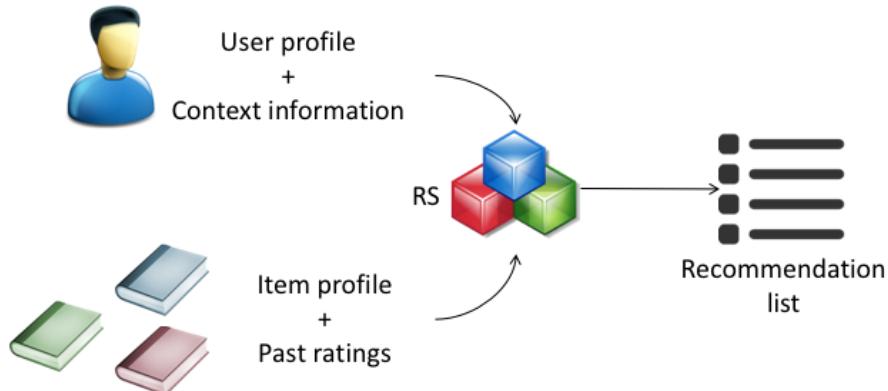
©2018, Karl Aberer, EPFL-IC, Laboratoire de systèmes d'informations répartis

Applied Classification - 4

In the collaborative paradigm, the recommender system uses the user profile (and possibly other information about the specific context in which the user needs a recommendation) and the data provided by other users to return a recommendation list with the most relevant items and their scores, from the highest (i.e., most relevant) to the lowest (i.e., least relevant). It is called collaborative because all the user participates to the recommendation by providing ratings to the items that they viewed/purchased etc.

Content-based Recommender System

Content-based = “Show me more of what I liked”



©2018, Karl Aberer, EPFL-IC, Laboratoire de systèmes d'informations répartis

Applied Classification - 5

In the content-based paradigm, the recommender system uses the profile of the items that the user rated in the past to return a recommendation list with the most relevant items and their scores, from the highest (i.e., most relevant) to the lowest (i.e., least relevant). It is called content-based because the items are not only objects, but they are defined by a set of attributes that summarize the content of the item.

There are also other approaches, such as hybrid ones, which combine collaborative and content-based, or knowledge-based ones, which require domain knowledge engineering to model how certain item features meet the user needs.

1. Collaborative Filtering

A widely used approach to generate recommendations

- Used by large e-commerce sites
- Applicable in many domains (e.g., books, movies ...)
- Well understood and studied

Approach (*Wisdom of the crowd*)

- Users give ratings to items
- Users with similar tastes in the past will have similar tastes in the future

Collaborative recommender systems are based on collaborative filtering. In collaborative filtering, given a user and an item not yet rated by the user, the goal is to estimate the user rating for this item by looking at the ratings for the same item that were given in the past by similar users. Collaborative filtering requires a community of users that provide ratings and a way to assess user similarity.

User-based Collaborative Filtering

Basic technique:

Given a user x and an item i not rated by x , estimate the rating $r_x(i)$ by

1. Find a set of users $N_U(x)$ who rated the same items as x (= neighbours of x) in the past and who have rated i
2. Aggregate the ratings of i provided by $N_U(x)$

Compute the ratings for all the items not rated by x and recommend the best-rated ones

The basic technique for user-based collaborative filtering needs to define

- 1 A metric to compute similarity between users
- 2 The number of neighbours considered for $N_U(x)$
- 3 A way to aggregate the ratings of i provided by $N_U(x)$

Similarity Between Users

Pearson correlation coefficient

$$sim(x, y) = \frac{\sum_{i=1}^N (r_x(i) - \bar{r}_x)(r_y(i) - \bar{r}_y)}{\sqrt{\sum_{i=1}^N (r_x(i) - \bar{r}_x)^2} \sqrt{\sum_{i=1}^N (r_y(i) - \bar{r}_y)^2}}$$

Cosine Similarity

$$sim(x, y) = \cos(\vartheta) = \frac{\sum_{i=1}^N r_x(i) \cdot r_y(i)}{\sqrt{\sum_{i=1}^N r_x(i)^2} \sqrt{\sum_{i=1}^N r_y(i)^2}}$$

x, y : users

N : number of items i rated by both x and y

$r_x(i)$: rating of user x of item i

\bar{r}_x : average ratings of user x

©2018, Karl Aberer, EPFL-IC, Laboratoire de systèmes d'informations répartis

Applied Classification - 8

A widely used similarity measure is the Pearson correlation coefficient, which returns a value between -1 and 1. The cosine similarity considers the users x and y as two N -dimensional vectors. Computing similarity follows exactly the same principles as for document similarity in vector space retrieval. If the angle between the two vectors is 0, thus the users are identical, and the cosine similarity returns 1. The greater the angle, the less similar the two users are. Assuming that all ratings are positive, the maximum angle is $\pi/2$, thus the minimum similarity is $\cos(\pi/2)=0$.

Example

	Item 1	Item 2	Item 3	Item 4	Item 5
U	5	3	4	4	???
User 1	3	1	2	3	3
User 2	4	3	4	3	5
User 3	3	3	1	5	4
User 4	1	5	5	2	1

$$\text{sim}_{\text{corr}}(U, \text{User1}) = 0.85$$

$$\text{sim}_{\text{cos}}(U, \text{User1}) = 0.97$$

$$\text{sim}_{\text{corr}}(U, \text{User2}) = 0.71$$

$$\text{sim}_{\text{cos}}(U, \text{User2}) = 0.99$$

$$\text{sim}_{\text{corr}}(U, \text{User3}) = 0$$

$$\text{sim}_{\text{cos}}(U, \text{User3}) = 0.89$$

$$\text{sim}_{\text{corr}}(U, \text{User4}) = -0.79$$

$$\text{sim}_{\text{cos}}(U, \text{User4}) = 0.79$$

This is an example of similarity computed using the two metrics. It is possible to see how the users are not ranked in the same way. For example, the most similar user to user U is User 1, if the Pearson correlation coefficient is used, or User 2, if the cosine similarity is used. Also User 4 is not that different from U using the cosine similarity, although he seems to have quite opposite tastes.

One drawback of Pearson correlation coefficient is that it is not computable, if the variance of one of the user ratings is 0 (e.g., a user with ratings 1 1 1 1).

However, in general, the correlation coefficient works well in general domains, particularly compared to the cosine similarity. One problem of cosine similarity is that it does not consider the absolute values of the ratings but only their relative distribution. This is appropriate when assessing the similarity of documents in information retrieval, since the absolute length of a document should not affect the similarity. However this does not work well with ratings. For example, two users with ratings 5 5 5 5 and 1 1 1 1 for item 1 to 4 have a cosine similarity of 1, since the vectors are parallel (albeit they have different magnitude). However, the users are not really similar, as the first likes everything while the second likes nothing.

Aggregate the Ratings

A common aggregation function is

$$r_x(a) = \bar{r}_x + \frac{\sum_{y \in N_U(x)} sim(x, y)(r_y(a) - \bar{r}_y)}{\sum_{y \in N_U(x)} |sim(x, y)|}$$

$N_U(x)$: neighbours of user x

a : item not rated by x

The aggregation function calculates whether the neighbours' ratings for the unseen item a are higher or lower than their average. We can consider this term as the bias of the user y w.r.t. item a . Then the function combines the rating bias using the similarity as a weight, so that the most similar neighbours will have more importance. Then the aggregated neighbours' bias is added/subtracted from user's x average rating.

Example

	Item 1	Item 2	Item 3	Item 4	Item 5
U	5	3	4	4	???
User 1	3	1	2	3	3
User 2	4	3	4	3	5
User 3	3	3	1	5	4
User 4	1	5	5	2	1

Closest users to U: $\text{sim}_{\text{corr}}(U, \text{User1}) = 0.85, \text{sim}_{\text{corr}}(U, \text{User2}) = 0.71$

$$\bar{r}_U = 4$$

$$(r_{U1}(I5) - \bar{r}_{U1}) = 3 - \frac{12}{5} = \frac{3}{5}, (r_{U2}(I5) - \bar{r}_{U2}) = 5 - \frac{19}{5} = \frac{6}{5}$$

$$r_U(I5) = 4 + \frac{0.85 * \frac{3}{5} + 0.71 * \frac{6}{5}}{0.85 + 0.71} = 4.87 \approx 5$$

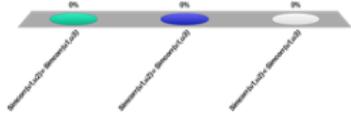
Given 3 users with ratings...

U1: 1 3

U2: 2 4

U3: 1 4

- A. $\text{Sim}_{\text{corr}}(u_1, u_2) > \text{Sim}_{\text{corr}}(u_1, u_3)$
- B. $\text{Sim}_{\text{corr}}(u_1, u_2) = \text{Sim}_{\text{corr}}(u_1, u_3)$
- C. $\text{Sim}_{\text{corr}}(u_1, u_2) < \text{Sim}_{\text{corr}}(u_1, u_3)$



User-based Collaborative Filtering

Problems

- Cold start: users or items without ratings
- Scalability: large numbers of users
- Data dispersion: highly variable ratings, difficult to find similar users

Possible solution

- Item-based collaborative filtering

A typical problem of user-based collaborative filtering is the cold-start problem, that is, the recommendation for a new user that did not rate anything yet or for an item that was never rated. Also, user-based collaborative filtering is problematic when there are millions of users and/or items (Amazon, Netflix). For example, for each user the most similar users have to be found (nearest neighbours) from a large set of users. With large numbers of users, it becomes hard to find common ratings between users, due to the high dispersion of data.

Item-based Collaborative Filtering

The basic technique: use the similarity between items (and not users) to make predictions

- Given a user x and an item i not rated by x , estimate the rating $r_x(i)$ by
 1. Find a set of items $N_I(i)$ which are similar to i and that have been rated by x
 2. Aggregate the ratings of the items in $N_I(i)$ and use the aggregation as prediction of $r_x(i)$

The basic technique for item-based collaborative filtering needs to define

- 1 A metric to compute similarity between items
- 2 The number of neighbours considered in N_I
- 3 A way to aggregate the ratings of the items in N_I

Similarity Between Items

	Item 1	Item 2	Item 3	Item 4	Item 5
U	5	3	4	4	???
User 1	3	1	2	3	3
User 2	4	3	4	3	5
User 3	3	3	1	5	4
User 4	1	5	5	2	1

$$\text{sim}_{\text{corr}}(\text{item5}, \text{item1}) = 0.97$$

$$\text{sim}_{\text{corr}}(\text{item5}, \text{item2}) = -0.48$$

$$\text{sim}_{\text{corr}}(\text{item5}, \text{item3}) = -0.43$$

$$\text{sim}_{\text{corr}}(\text{item5}, \text{item4}) = 0.58$$

The attributes that define an item are the ratings of the users different than U. The similarity can be computed with the same functions explained before, correlation coefficient or cosine similarity. In this example, similarity is computed with the correlation coefficient, and it is possible to see how item 1 and 4 are the most similar to item 5. Thus, aggregating the ratings of these items given by user U is an estimate of the rating of user U for item 5.

Aggregate the Ratings

A common aggregation function is

$$r_x(a) = \frac{\sum_{b \in N_I(a)} sim(a, b)r_x(b)}{\sum_{b \in N_I(a)} |sim(a, b)|}$$

$N_I(a)$: neighbours of item a

b : item rated by x

The aggregation function computes the prediction of an item a for a user x by computing the sum of the ratings given by the user on the items similar to a . Each rating is weighted by the corresponding similarity $sim(a,b)$ between items a and b .

Example

	Item 1	Item 2	Item 3	Item 4	Item 5
U	5	3	4	4	???
User 1	3	1	2	3	3
User 2	4	3	4	3	5
User 3	3	3	1	5	4
User 4	1	5	5	2	1

Closest items: $\text{sim}_{\text{corr}}(\text{item5}, \text{item1}) = 0.97$, $\text{sim}_{\text{corr}}(\text{item5}, \text{item4}) = 0.58$

$$r_U(\text{item5}) = (0.97*5 + 0.58*4)/(0.97 + 0.58) = 4.62 \approx 5$$

Item-based Collaborative Filtering

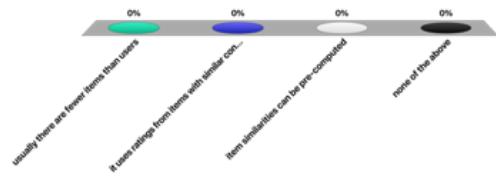
Item-based collaborative filtering does not solve the cold-start problem but has better scalability

- Calculate all the pair-wise item similarities in advance
- Items similarities are more stable
- The neighbourhood $N_I(a)$ to be considered at runtime is small

Although item-based collaborative filtering does not solve the cold-start problem, it is usually more suited to tackle big datasets. In fact, the item similarities can be computed in advance, given that the items are more stable than the users (in general, new items appear at a slower pace than new users). Furthermore, the size of the neighbourhood of a given item a is in general smaller than the neighbourhood of a user x , because $N_I(a)$ is composed of the other items rated by user x (e.g., tens of DVDs bought on Amazon), while $N_U(x)$ is composed by the other users that rated the same items as x (e.g., millions of users that watched The Godfather on Netflix).

Item-based collaborative filtering addresses better the cold-start problem because ...

- A. usually there are fewer items than users
- B. it uses ratings from items with similar content
- C. item similarities can be pre-computed
- D. none of the above



2. Content-based Recommendation

Collaborative filtering needs only the ratings, it does not require any information about the items

However, the *content* of the item might provide some useful information

- E.g., recommend sci-fi novels to users who liked Asimov books in the past

Both recommendation techniques we saw before need only the ratings of the user, and no information about the item is needed. However, the content of the item might provide useful information for the recommendation. Content-based recommendation takes into consideration only the items that have been rated by the user and the content of all existing items, significantly reducing the scalability problems since a community of users is not necessary any more.

Content-based Recommendation

The basic technique

- Given the items that have been rated by user x in the past
 1. Find the items that are similar to the past items
 2. Aggregate the ratings of the most similar items

The basic technique for content-based recommendation needs to define

- 1 A way to formalize the item content
- 2 A similarity measure between items
- 3 An aggregation function for the ratings

Content-based Recommendation

Most methods originate from information retrieval to extract the content of an item

Given an item and a textual description (e.g. movie synopsis, book review), compute the tf-idf weights

$$w(t, a) = tf(t, a) \cdot idf(t) = \frac{freq(t, a)}{\max_{s \in T} freq(s, a)} \cdot \log\left(\frac{N}{n(t)}\right)$$

N : number of items to recommend

$n(t)$: number of items where term t appears

The first step for content-based recommendation is the characterisation of the items in terms of TF-IDF. Each item can be therefore described by the set of terms that appear in their description, with their associated weight (the TF-IDF measure)

To compute the TF-IDF measure, all the steps typical of information retrieval (removing stopwords, stemming, only top M terms etc.) are applied.

Similarity between Items

Cosine similarity

$$sim(a,b) = \cos(\vartheta) = \frac{\sum_{t=1}^T w(t,a) \cdot w(t,b)}{\sqrt{\sum_{t=1}^T w(t,a)^2} \sqrt{\sum_{t=1}^T w(t,b)^2}}$$

a, b : items

T : number of terms t that appear in all items

$w(t, a)$: tf-idf of term t in item a

The cosine similarity in content-based recommendation considers the items a and b as two T -dimensional vectors, where each dimension is the TF-IDF measure of a specific term t .

Making Predictions

Given a set of items D that have been rated by the user,
find the nearest neighbours $N_I(a)$ in D of a new item a

Use the ratings of the nearest neighbours to predict the rating of a with the aggregation function seen before

$$r_x(a) = \frac{\sum_{b \in N_I(a)} sim(a, b)r_x(b)}{\sum_{b \in N_I(a)} |sim(a, b)|}$$

Once a way to compute the similarity between items has been established, using the vector space model (TF-IDF + cosine similarity), the simplest approach to estimate the rating of item a not yet seen/rated by user x is to find the nearest neighbours of a in the subset of items that have been already rated by the user x and aggregate their ratings. Note that this approach to predict the rating is analogous to the use of the kNN nearest neighbor classification for document classification. In both cases, the nearest documents are used as predictors for the most likely label respectively rating.

Content-based Recommendation

Problems

- Cold start problem
- Content can be limited or impossible to extract (multimedia)
- Tends to recommend “more of the same”

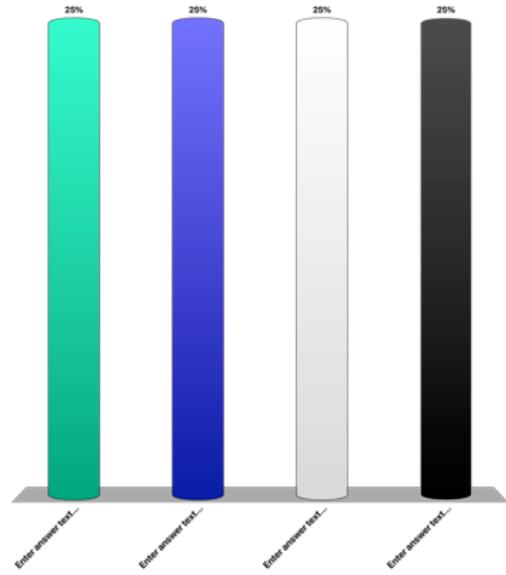
More scalable: tf-idf of items can be computed offline

Content-based recommendation also suffers from the cold start problem as collaborative filtering, but only for the users that did not rate anything in the past. Another problem is that sometimes the information available to extract the content is limited (e.g., movie synopsis versus full plot) or impossible to process (e.g. a sound or a video). Finally, content-based recommendation tends to recommend more of the same, it does not surprise the user with new interesting items.

On the other hand, it is in general more scalable, because it does not rely on a community of users to provide recommendations. Furthermore, the TF-IDF measure can be computed once a new item enters the system, and then update the pair-wise similarities with all the existing items.

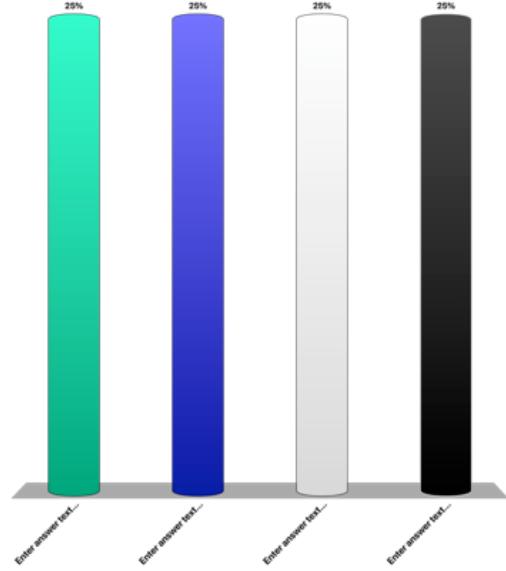
For a user that has not done any ratings, which method can make a prediction?

- A. User-based collaborative RS
- B. Item-based collaborative RS
- C. Content-based RS
- D. None of the above



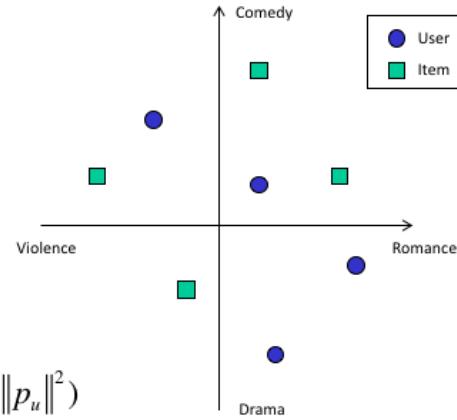
For an item that has not received any ratings, which method can make a prediction?

- A. User-based collaborative RS
- B. Item-based collaborative RS
- C. Content-based RS
- D. None of the above



Advanced Methods: Matrix Factorization

	I1	I2	I3	I4
U1	3	1	-	-
U2	-	3	4	3
U3	3	-	1	-
U4	-	5	5	2



$$\min_{q,p} \sum_{(u,i) \in M} (r_u(i) - q_i^T p_u)^2 + \lambda (\|q_i\|^2 + \|p_u\|^2)$$

©2018, Karl Aberer, EPFL-IC, Laboratoire de systèmes d'informations répartis

Applied Classification - 28

Matrix factorization transforms the user-item rating matrix into a latent factor model, where each user and item is described as a vector in a d-dimensional space. It can be understood as a combination of user-based and item-based collaborative filtering.

For the items, the factors describe evident as well as non interpretable characteristics. For the users, each factor measures how much the user likes items that score high on the corresponding factor. The estimate of the rating of item a that user u would give can be computed as the dot product between the two vectors: $q_i^T \cdot p_u$ where q_i is the item vector and p_u is the user vector. The major challenge of course is computing the mapping of each item and user to their corresponding vectors q_i and p_u . After that, the rating can be estimated using the dot product.

This approach is analogous to singular value decomposition for document retrieval. However, a direct application of SVD is not possible because the user-item matrix is not complete (missing ratings). Thus, to learn the factor vectors q_i and p_u the system minimises the regularised square error on the set of known ratings $(u,i) \in M$. The minimization problem can be solved for instance using stochastic gradient descent.

8. MINING SOCIAL GRAPHS

Analyzing Graphs

Up to now

- Objects have been described by attributes
- Relationships among objects inferred from distance measure on attributes

Use of explicit relationships: Graphs

- In many cases relationships are explicitly given
- Prime example: social network graphs
- Graph structure can be inferred from distance measure (e.g. for documents)

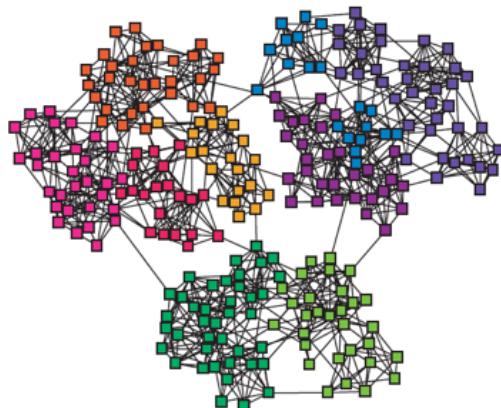
So far we have assumed that relationships among different objects in clustering are implicitly defined, using a distance measure that is based on the objects' attributes. In many applications such relationships are not implicitly given, but explicitly provided, notably in social network graphs. There the relationships among users are given by different friend relationships or interactions such as likes and retweets. The resulting graphs can be weighted or unweighted. In the following we will consider the case of unweighted graphs.

Of course, it would also be possible to derive a relationship graph by using a distance measure based on attribute values and using the distance as an edge weight, respectively consider only edges above a distance threshold. In this way the following methods of clustering objects based on graphs could also be applied in this more traditional context.

Graphs and Clustering

Networks often contain structure

- Clusters (also called communities, modules)



©2018, Karl Aberer, EPFL-IC, Laboratoire de systèmes d'informations répartis

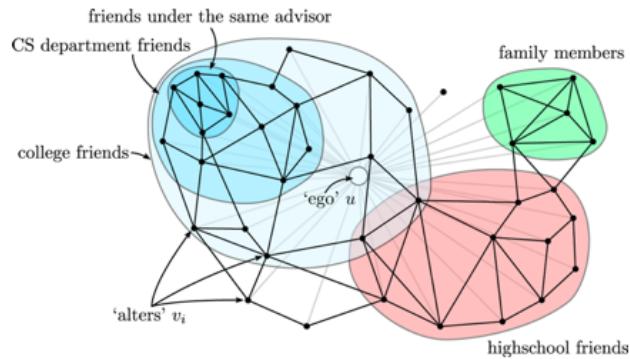
Applied Classification - 31

It is widely observed that natural networks contain structure. This is true for social networks (e.g. on social media platforms, citation networks), as well as for many natural networks as we find them in biology. Graph-based clustering aims at uncovering such hidden structures.

Social Network Analysis

Clusters (communities) in social networks
(Twitter, Facebook) related

- Interests
- Level of trust

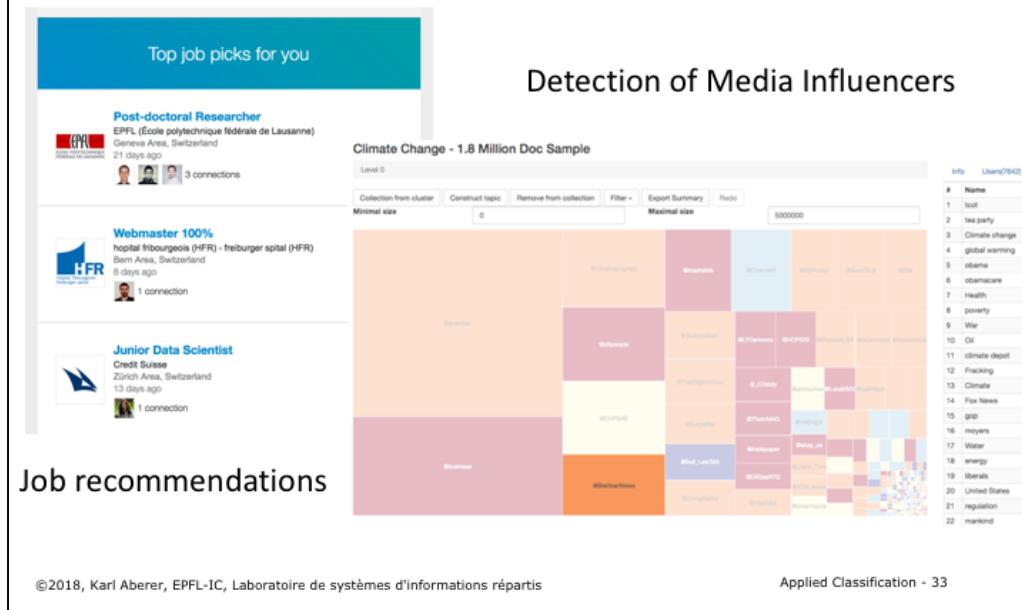


©2018, Karl Aberer, EPFL-IC, Laboratoire de systèmes d'informations répartis

Applied Classification - 32

In social networks mining community structures is particularly popular. Consider the social network neighborhood of a particular social network user, i.e. all other social network accounts that are connected to the user, either through explicit relationships, such as a follower graph, or through interactions such as likes or retweets. Typically the social neighborhood would decompose into different groups, depending on interests and social contexts. For example, the friends from high school would have a much higher propensity to connect to each other, than with members of the family of the user. Thus they are likely to form a cluster. Similarly, other groups with shared interest or high mutual level of trust would form communities.

Use of Community Structures



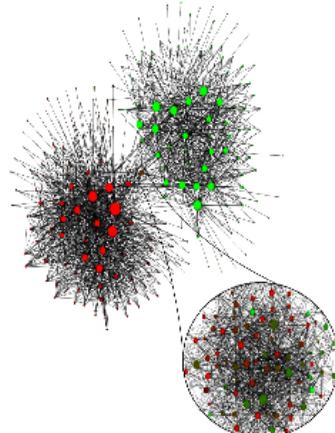
Many tasks can benefit from a reliable community detection algorithm. Online Social Networks rely on their underlying graph to recommend content, for example relevant jobs. Knowing to which community a user belongs can improve dramatically the quality of such recommendations.

Another typical use of community detection is to identify media influencers. Community detection can first be used to detect communities that share in social networks common interests or beliefs (e.g. for climate change we might easily distinguish communities that are climate deniers and climate change believers), and then the main influencers of such communities could be identified.

Use of Community Structures: Social Science

Call patterns in Belgium mobile phone network

- Two almost separate communities



V.D. Blondel et al, *J. Stat. Mech.* P10008 (2008).

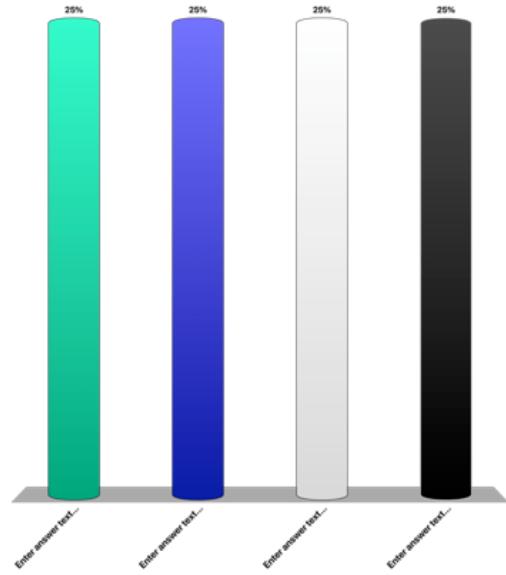
©2018, Karl Aberer, EPFL-IC, Laboratoire de systèmes d'informations répartis

Applied Classification - 34

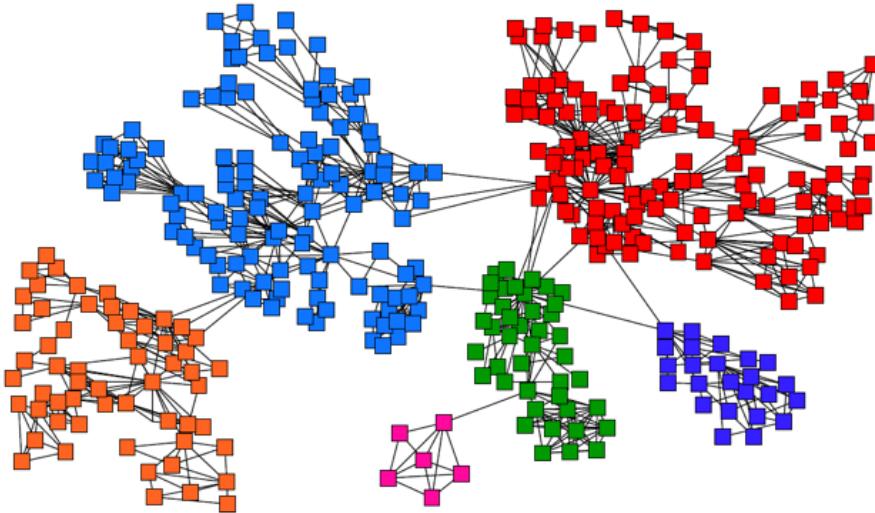
In 2008 Vincent Blondel and his students have started applying a new community detection algorithm to the call patterns of one of the largest mobile phone operators in Belgium. It was designed to identify groups of individuals who regularly talk with *each other* on the phone, breaking the whole country into numerous small and not so small communities by placing individuals next to their friends, family members, colleagues, neighbors, everyone whom they regularly called on their mobile phone. The result was somewhat unexpected: it indicated that Belgium is broken into two huge communities, each consisting of many smaller circles of friends. Within each of these two groups the communities had multiple links to each other. Yet, these communities never talked with the communities in the other group (guess why?). Between these two mega-groups was sandwiched in a third, much smaller group of communities, apparently mediating between the two parts of Belgium.

Which of the following techniques do you believe would be most appropriate to identify communities on a social graph?

- A. Cliques
- B. Random Walks
- C. Shortest Paths
- D. Association rules



Find Densely Linked Clusters



©2018, Karl Aberer, EPFL-IC, Laboratoire de systèmes d'informations répartis

Applied Classification - 36

The intuition behind community detection is that the heavily linked components of the graph belong to the same community.

Similarly as earlier for clustering in multidimensional spaces, therefore, the goal of a community detection algorithm is to identify communities that are heavily intra-linked (high intra-cluster similarity) and scarcely inter-linked (low inter-cluster similarity).

Types of Community Detection Algorithms

Hierarchical clustering

- iteratively identifies groups of nodes with high similarity

Two strategies

- **Agglomerative algorithms** merge nodes and communities with high similarity
- **Divisive algorithms** split communities by removing links that connect nodes with low similarity

In general, community detection algorithms are based on a hierarchical approach. The idea is that within communities sub-communities can be identified, till the network decomposes into individual nodes. In order to produce such a hierarchical clustering, two approaches are possible: either by starting from individual nodes, by merging them into communities, and recursively merge communities into larger communities till no new communities can be formed (agglomerative algorithms), or by decomposing the network into communities, and recursively decompose communities till only individual nodes are left (divisive algorithms).

In the following we will present one representative of each of the two categories of algorihtms:

1. The Louvain Algorithm, an agglomerative algorithm
2. The Girvan-newman algorithm, a divisive algorithm

Louvain Modularity Algorithm

Agglomerative Community Detection

- Based on a measure for community quality (**Modularity**)
- greedy optimization of modularity

Overall algorithm

- **first** small communities are found by optimizing modularity **locally** on all nodes
- then each small community is **grouped** into one new community node
- **Repeat** till no more new communities are formed

©2018, Karl Aberer, EPFL-IC, Laboratoire de systèmes d'informations répartis

Applied Classification - 38

The Louvain algorithm is essentially based on the use of a measure, modularity, that allows to assess the quality of a community clustering. The algorithm performs greedy optimization of this measure. It is fairly straightforward: initially every node is considered as a community. The communities are traversed, and for each community it is tested whether by joining it to a neighboring community, the modularity of the clustering can be improved. This process is repeated till no new communities form anymore.

Measuring Community Quality

Communities are sets of nodes with many mutual connections, and much less connections to the outside

Modularity measures this quality: the higher the better

$$\sum_{c \in \text{Communities}} (\#\text{edges within } c - \text{expected } \#\text{edges within } c)$$

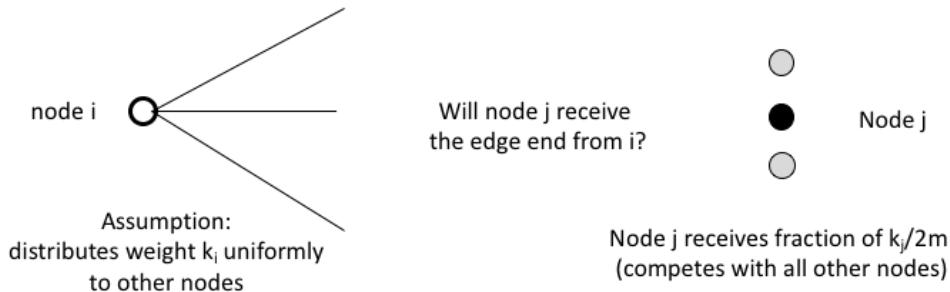
In order to measure the quality of a community clustering, modularity compares simply the difference between the number of edges that occur within a community with the number of edges that would be expected if the edges in the graph would occur randomly. The random graph model is used as what is called the null model, a graph that has the same distribution of node degrees, but randomly assigned connections.

Expected Number of Edges

Graph with unweighted edges

- m = total number of edges
- k_i = number of outgoing edges of node i (degree)

Observation: edges connect to $2m$ nodes



©2018, Karl Aberer, EPFL-IC, Laboratoire de systèmes d'informations répartis

Applied Classification - 40

The null model requires to determine the number of edges that we would expect among a set of nodes, of which we know the degrees, but not the connectivity in detail. Assume that for all nodes i in a community we would now their degree k_i (number of edges leaving the node). How many edges would we then observe if the connections on the graph were generated randomly? To answer this question we reason as follows: select one of the nodes i with degree k_i . What is now the probability (or fraction of weight) an arbitrary other node j with weight k_j would receive? If there are a total of m edges in the network, there are $2m$ edge ends (since each edge ends in two nodes). If the edge ends of k_i are uniformly distributed, node j would thus receive $k_i/2m$ edge ends. Thus the number of expected edges connecting nodes i and j would be k_i times $k_j/2m$.

Modularity

We define now the modularity measure Q

- A_{ij} = effective number of edges between nodes i and j
- c_i, c_j = communities of nodes i and j

Effective number of edges Expected number of edges

$$Q = \frac{1}{2m} \sum_{i,j} \left(A_{ij} - \frac{k_i k_j}{2m} \right) \delta(c_i, c_j)$$

Properties

- Q in $[-1,1]$
- $0.3-0.7 < Q$ means significant community structure

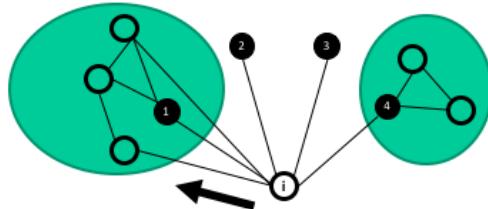
Given the expected number of edges (the null model) we can now formulate the modularity measure as the total difference of number of expected and effective edges for all pairs of nodes from the same cluster. The delta function assures that only nodes belonging to the same community are considered, since it returns 1 when $c_i = c_j$.

Due to normalization the measure returns values between -1 and 1. In general, if modularity exceeds a certain threshold (0.3 to 0.7) the clustering of the network is considered to exhibit a good community structure.

Locally Optimizing Communities

What is the modularity gain by moving node i to the communities of any of its neighbors?

- Test all possibilities and choose the best

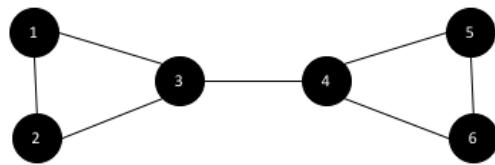


Given the modularity measure the next question is now how to use it to infer the communities. This is performed through local optimization. The algorithm sequentially traverses all nodes of the network, and for each node checks how the modularity can be increased maximally by having the node joining the node of a neighboring community. It then decides to join that node to the best community.

Example

Initial modularity $Q = 0$

Start processing nodes in order



We illustrate the algorithm for a simple example. Initially, the modularity is zero since all nodes belong to different communities. We start now to process the nodes in some given order, e.g. size of identifier.

Example: Processing Node 1

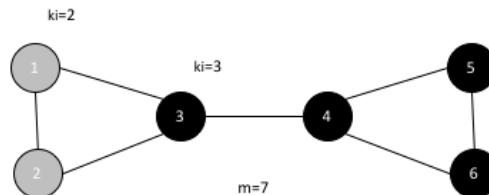
Joining node 1 to node 2

- $Q = 1/2 * 7(1 - 2 * 2/2 * 7) = 1/14 (1 - 4/14) = 1/14 * 10/14 > 0$

Joining node 1 to node 3

- $Q = 1/14 (1 - 2 * 3/2 * 7) = 1/14 (1 - 6/14) = 1/14 * 8/14 > 0$

New modularity: $1/14 * 10/14$



For node 1 we can join it either to node 2 or 3, it's two neighbors. To decide which is better, we have to compute modularity after the join. We see that it is better to join node 2, as the resulting modularity will be higher. We can think of this as follows: since node 3 has more connections, it is more likely to be randomly connected to node 1, this connecting to node 2 is more "surprising".

Example: Processing Nodes 2 and 3

Joining node 2 to node 3 (leaving community of Node 1)

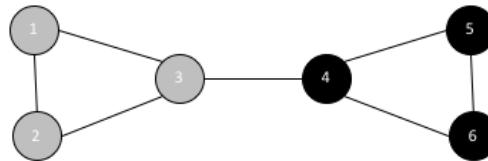
- no improvement

Joining node 3 to community {1,2} (via node 1 or 2)

- $Q = 1/14 (3 - 4/14 - 6/14 - 6/14) = 1/14 * 26 / 14$

Joining node 3 to node 4

- $Q = 1/14 10/14 + 1/14 (1 - 9/14) = 1/14 * 15/14$



For node 2 there will be no change, as the only alternative would be node 3. But for the same reasons node 1 did not join node 3, also node 2 does not. Next is node 3: here we have two choices, either to join community {1,2}, or to join node 4. In the first case we obtain a larger community, and in the second case two smaller communities. Computation of modularity reveals that joining 1 and 2 gives a much better community structure.

Example: Processing Nodes 4, 5 and 6

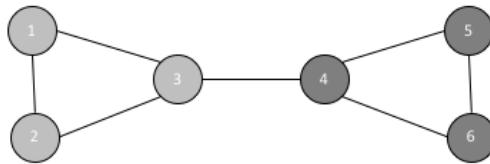
Joining node 4 to community {1,2,3} via 3

- $Q = 1/14 (4 - 4/14 - 6/14 - 6/14 - 6/14 - 6/14 - 9/14) = 1/14 * 19/14$

Joining node 4 to node 5

- $Q = 1/14 * 26/14 + 1/14 * (1 - 6/14) = 1/14 * 34/14$

Finally, also node 6 will join the second community



Similar arguments as before will then join node 4 to 5 and finally node 6 to the community consisting of nodes 4 and 5. Then the first round of processing is over.

Example: Merging Nodes

Now that all nodes have been processed, we merge nodes of the same community in a single new nodes and restart processing

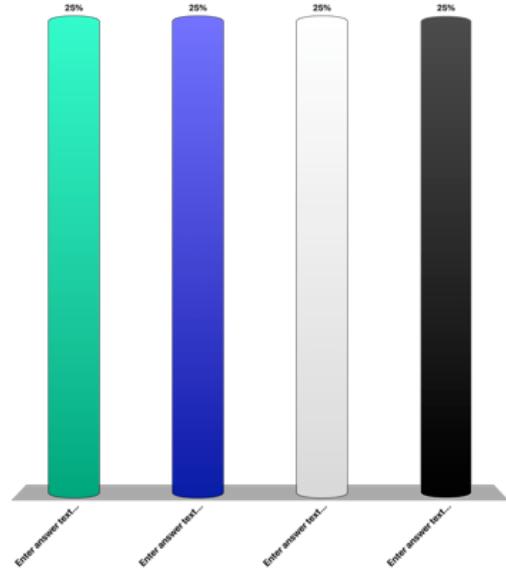
Will the two remaining nodes merge? Answer: yes



For the next round of processing the resulting community nodes are collapsed in new nodes for the complete community, and the algorithm is re-run with the new resulting graph. Obviously, now the two remaining nodes will merge into a community, as the modularity will move from zero to positive. Then the algorithm terminates.

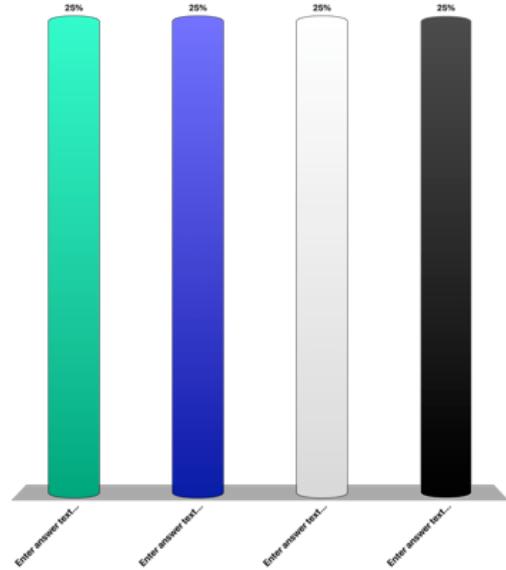
Modularity clustering will end up always with a single community at the top level?

- A. true
- B. Only for dense graphs
- C. Only for connected graphs
- D. never



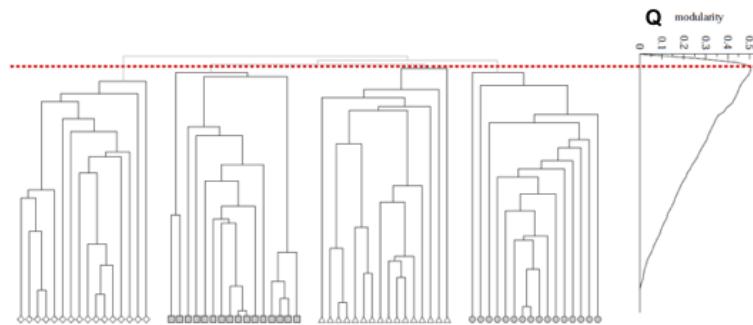
Modularity clustering will end up always with the same community structure?

- A. true
- B. Only for connected graphs
- C. Only for cliques
- D. false



Modularity to Evaluate Community Quality

Modularity can also be used to evaluate the best level to cutoff of a hierarchical clustering



©2018, Karl Aberer, EPFL-IC, Laboratoire de systèmes d'informations répartis

Applied Classification - 50

Apart from constructing the communities, the modularity measure can also be used to evaluate the quality of communities in hierarchical clustering. This can be done independently of how the clustering has been constructed. In fact, there is an optimal level of clustering when moving from one level of the hierarchy to the next. Initially increasing the number of communities increases their quality, by separating distinct communities. At a certain point, splitting of communities in smaller communities worsens the quality of the community structure. Thus there exists an optimum point of clustering that can be selected using modularity.

Louvain Modularity - Discussion

Widely used in social network analysis and beyond

- Method to extract communities from very large networks very fast

Complexity: $O(n \log n)$

Louvain modularity clustering is today the method of choice for social network clustering, mainly because of its good computational efficiency. It runs in $n \log n$, which makes it applicable for very large networks, as they occur today, in particular for social networks resulting from large platforms, as social network sites, messaging services or telephony.

Girvan-Newman Algorithm

Divisive Community Detection

- Based on a **betweenness measure** for edges, measuring how well they separate communities
- Decomposition of network by splitting along edges with highest separation capacity

Overall algorithm

- Repeat until no edges are left
 - Calculate betweenness of edges
 - Remove edges with highest betweenness
 - Resulting connected components are communities
- Results in hierarchical decomposition of network

We now introduce a second algorithm for community detection, that belongs to the class of divisive algorithms. Also this algorithm is based on a measure, this time on a measure on edges. The betweenness measure gives an indication which edges are likely to connect different communities, and thus are good splitting points, to partition larger parts of the network into communities. The algorithm recursively decomposes the network, by removing edges with the highest betweenness measure, till no edges are left. Also this algorithm results in a hierarchical clustering.

Edge Betweenness

Edge betweenness: fraction of number of shortest paths passing over the edge

$$\text{betweenness}(v) = \sum_{x,y} \frac{\sigma_{xy}(v)}{\sigma_{xy}}$$

where

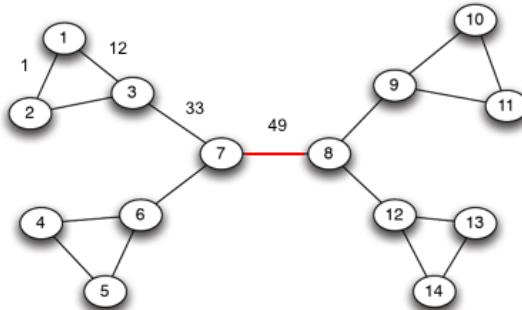
σ_{xy} : number of shortest paths from x to y

$\sigma_{xy}(v)$: number of shortest paths from x to y
passing through v

Betweenness centrality is an indicator of a node's centrality in a network. It is equal to the number of shortest paths from all vertices to all others that pass through that node. A node with high betweenness centrality has a large influence on the transfer of items through the network, under the assumption that item transfer follows the shortest paths. The concept finds wide application, including computer and social networks, biology, transport and scientific cooperation.

Alternatively, there exist also the concept of *random-walk betweenness*. A pair of nodes m and n are chosen at random. A walker starts at m , following each adjacent link with equal probability until it reaches n . Random walk betweenness x_{ij} is the probability that the link $i \rightarrow j$ was crossed by the walker after averaging over all possible choices for the starting nodes m and n .

Example

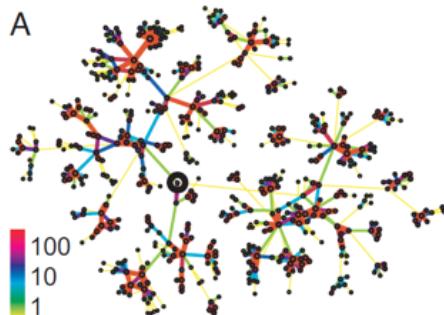


©2018, Karl Aberer, EPFL-IC, Laboratoire de systèmes d'informations répartis

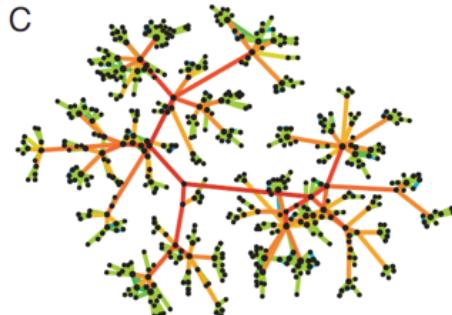
Applied Classification - 54

In this example network we compute betweenness for some selected edges. For example, for the edge 1-2 there is one shortest path between 1 and 2 that traverses the edge 1-2, thus the value is 1. For 1-3 there are shortest paths from the remaining 12 nodes in the network (except node 2) to node 1 that have to pass through this edge, thus the betweenness value is 12. For edge 3-4 we have paths going to both nodes 1 and 2, thus the betweenness value of that edge is significantly higher than for 1-3.

Underlying Intuition



**Edge strengths (call volume)
in a real network**



**Edge betweenness
in a real network**

©2018, Karl Aberer, EPFL-IC, Laboratoire de systèmes d'informations répartis

Applied Classification - 55

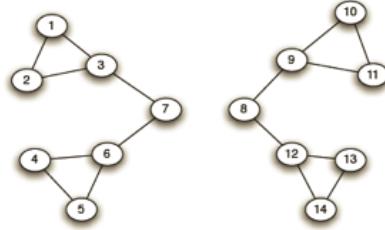
In a sense betweenness is the dual concept to connectivity. This is illustrated by the two graphs that result from real phone call networks. On the left hand side we see the indication of the strengths of connections among the nodes.

Communities are tightly connected by such links. On the right hand side we see the betweenness measure. Now the links that are connecting communities have a high strength, since, intuitively speaking, the traffic from one community to another has to traverse over these sparsely available links.

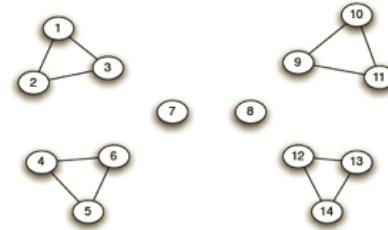
Example

Need to re-compute betweenness at every step

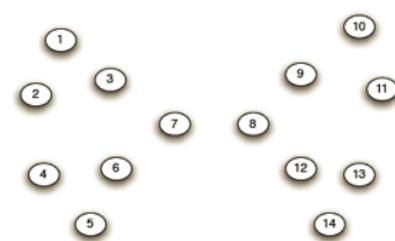
Step 1



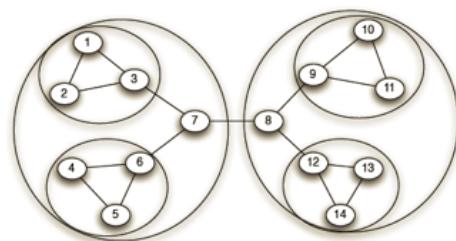
Step 2



Step 3



Hierarchical network decomposition



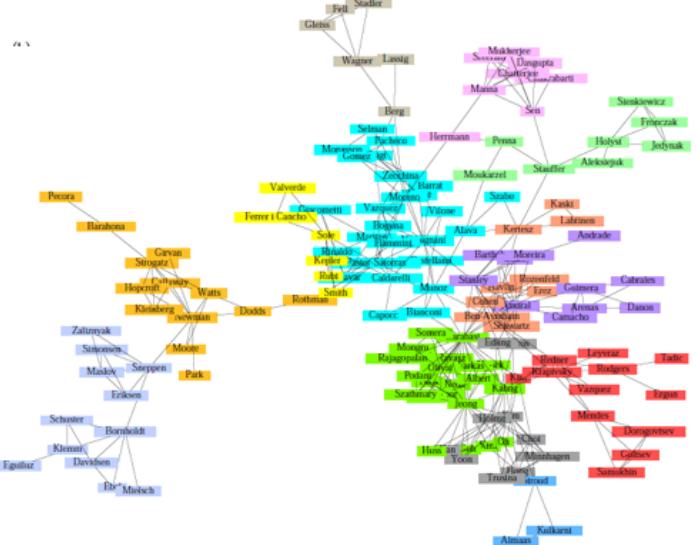
©2018, Karl Aberer, EPFL-IC, Laboratoire de systèmes d'informations répartis

Applied Classification - 56

Here we illustrate the execution of the Girvan-Newman algorithm. In Step 1 we remove one edge (in the middle) that had the highest betweenness value, resulting in two communities. Next (by symmetry) the edges connected to nodes 7 and 8 are removed, and in the third and fourth step the network decomposes completely. By overlaying the communities that have resulted from each step we obtain the final hierarchical clustering.

As the graph structure changes in every step, the betweenness values have to be recomputed in every step. This constitutes the main cost of the algorithm.

Girvan-Newman: Sample Results



Communities in physics collaborations

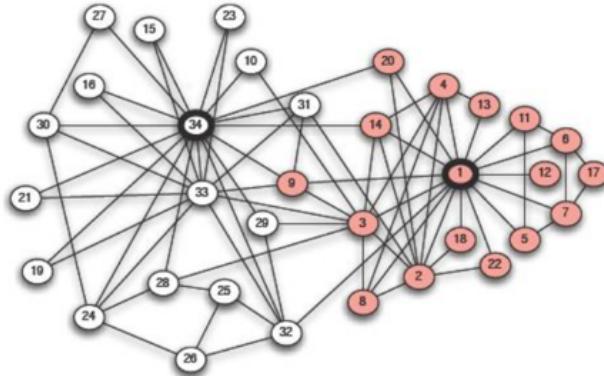
©2018, Karl Aberer, EPFL-IC, Laboratoire de systèmes d'informations répartis

Applied Classification - 57

The algorithm has been applied in many contexts, in particular on smaller graphs resulting from social science studies.

Girvan-Newman: Sample Results

Zachary's Karate club: Hierarchical decomposition



©2018, Karl Aberer, EPFL-IC, Laboratoire de systèmes d'informations répartis

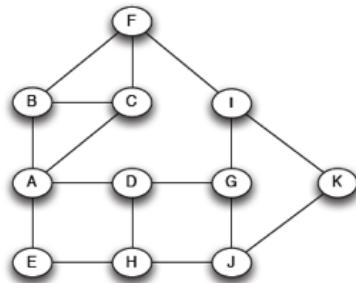
Applied Classification - 58

The origins of the algorithm come from studying a network of the 34 karate club members studied by the sociologist Wayne Zachary. Links capture interactions between the club members outside the club. The circles and the squares denote the two fractions that emerged after the club split into two following a feud between the group's president and the coach. The split between the members closely follows the boundaries of these communities. This karate club has been historically used as a benchmark to test community finding algorithms.

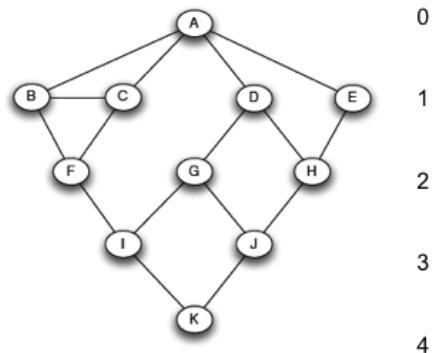
There was one outlier, node number 3, where the algorithm wrongly assigns the member. at the time of conflict, node 9 was completing a four-year quest to obtain a black belt, which he could only do with the instructor (node 34)

Computing Betweenness - BFS

Computing
betweenness of paths
starting at node A



Perform BFS
starting from A



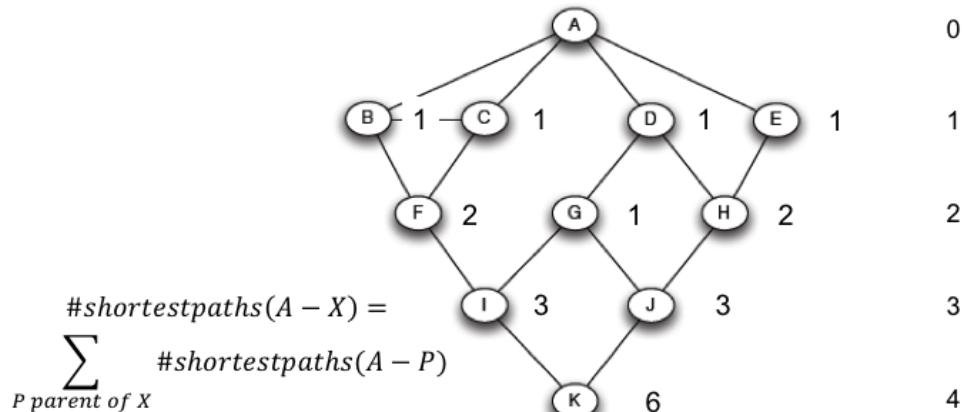
©2018, Karl Aberer, EPFL-IC, Laboratoire de systèmes d'informations répartis

Applied Classification - 59

We describe now the process of computing the betweenness values. The approach is to perform a breadth-first search (BFS) for every node in the graph. The nodes are arranged in increasing levels of distance of the starting node, e.g. node A.

Computing Betweenness – Path Counting

Count the number of shortest paths from A to all other nodes of the network



©2018, Karl Aberer, EPFL-IC, Laboratoire de systèmes d'informations répartis

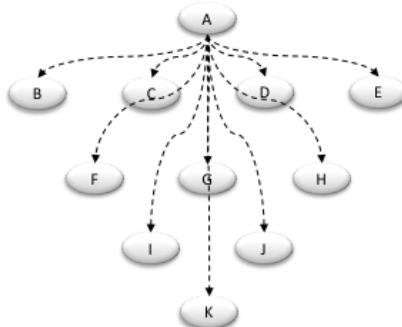
Applied Classification - 60

In a first phase we count the number of shortest paths that are leading to each node, starting from node A. To do so we can simply reuse the data that has been computed at the previous level, summing up the number of paths that have been leading to each parent of a given node.

Computing Betweenness – Edge Flow

Edge Flow

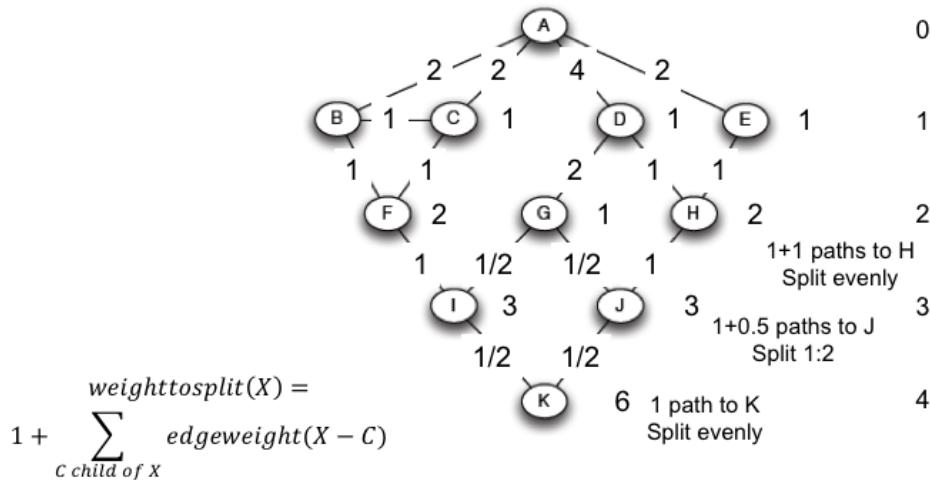
- 1 unit of flow from A to each node
- Flow to be distributed evenly over all paths
- Sum of the flows from all nodes equals the betweenness value



©2018, Karl Aberer, EPFL-IC, Laboratoire de systèmes d'informations répartis

Applied Classification - 61

Computing Edge Flow



©2018, Karl Aberer, EPFL-IC, Laboratoire de systèmes d'informations répartis

Applied Classification - 62

In a second phase we compute edge flow values (that are the betweenness values related to node A). We proceed in a bottom-up approach. The flow that arrives at every node is 1. In addition, the node receives also all the flows that are passing on to the children. Thus, the total flow weight a node receives (or has to pass through) is 1 plus the flows to all of its children. This weight is distributed over the parents, proportionally to the number of paths that are leading to those parents (what has been computed in phase 1).

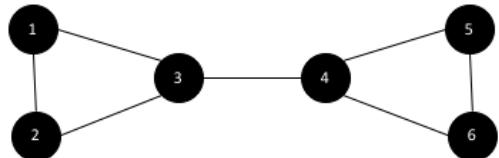
Algorithm for Computing Betweenness

1. Build one BFS structure for each node
2. Determine edge flow values for each edge using the previous procedure
3. Sum up the flow values of each edge in all BFS structures to obtain betweenness value
 - Flows are computed between each pairs of nodes
→ final values divided by 2

Once the flows specific to a node have been computed for every node, the last step is to aggregate for each edge all the flow values that have been computed for all the nodes. In this way we compute each flow twice (flow into both direction), therefore, the final betweenness value corresponds to the aggregate flow value divided by 2.

Betweenness of edge 3-4 is ...

- A. 16
- B. 12
- C. 9
- D. 4

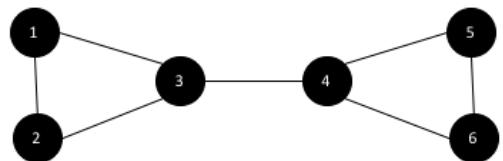


©2018, Karl Aberer, EPFL-IC, Laboratoire de systèmes d'informations répartis



When computing path counts for node 1 with BFS, the count at 6 is ...

- A. 1
- B. 2
- C. 3
- D. 4



©2018, Karl Aberer, EPFL-IC, Laboratoire de systèmes d'informations répartis



Applied Classification - 65

Girvan-Newman Discussion

Classical method

- Works for smaller networks

Complexity

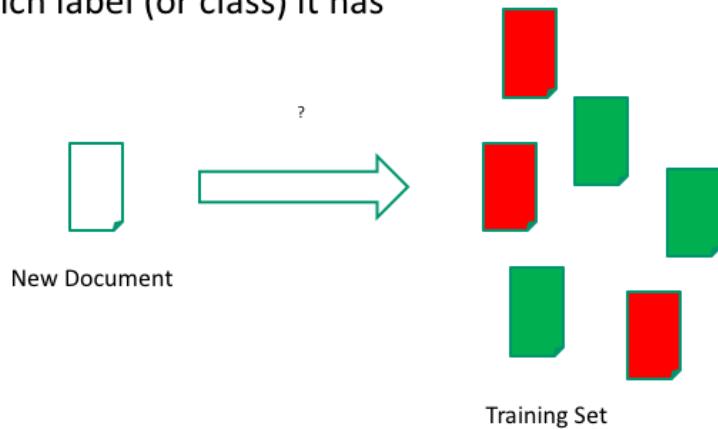
- Computation of betweenness for one link: $O(n^2)$
- Computation of betweenness for all links: $O(L n^2)$
- Sparse matrix: $O(n^3)$

The Girvan-Newman algorithm is the classical algorithm for community detection. Its major drawback is its scalability. The flow computation for one link has quadratic cost in the number of nodes (it has to be computed for each pair of nodes). If we assume sparse networks, where the number of links is of the same order as the number of nodes, the total cost is cubic. This was also one of the motivations that inspired the development of the modularity based community detection algorithm.

9. DOCUMENT CLASSIFICATION

Document Classification

Task: Given a training set of labelled documents, construct a **classifier** decide for an unlabeled document which label (or class) it has



Document classification is a special case of classification where the objects to be classified are (unstructured documents). This task has huge importance in many application areas, such as spam filtering, sentiment analysis, document filtering etc. and has therefore been studied as a specific classification problem in very much detail.

Document Classifiers

Features

- Words of the documents
 - bag of words
 - document vector
- More detailed information on words
 - Phrases
 - Word fragments
 - Grammatical features
- Any metadata known about the document and its author

An important question in document classification is the choice of features. A wide variety of possible features can be derived from text documents, many of which correspond to features used in document representation of information retrieval.

Example



PromotionDesSciences @EPFL_SPS · May 1

More than 4000 visitors attended Scientastic the science festival of EPFL,
organised in Valais for the first time. actu.epfl.ch/news/scientast...

Bag of words: {more, than, visitors, attend, ...}

Phrases: {"science festival", "first time"}

Word fragments: {mor, ore, tha, han, vis, isi, ..}

Grammatical features:

More than 4000 visitors attended Scientastic the science festival of EPFL , organised
in Valais for the first time

Adjective	Yellow
Adverb	Brown
Conjunction	Light Green
Determiner	Dark Blue
Noun	Green
Numeral	Blue
Preposition	Pink
Pronoun	Light Green
Verb	Yellow

<http://parts-of-speech.info/>

Author: [@EPFL SPS](https://twitter.com/EPFL_SPS)

©2018, Karl Aberer, EPFL-IC, Laboratoire de systèmes d'informations répartis

Applied Classification - 70

Here we give a few examples of features that could be extracted from a simple tweet, as shown above.

Challenge in Document Classification

The feature space is of very high dimension

- Vocabulary size
- All word bigrams
- All character trigrams
- etc.

Dealing with high dimensionality

- Feature selection, e.g., mutual information
- Dimensionality reduction, e.g., word embedding
- Classification algorithms that scale well

A problem that is characteristics for document classification is the fact that the feature space for documents can be huge. From information retrieval we know that the vocabulary size of a large document collection can grow to significant sizes, reaching millions of terms. When considering more features, such as word bigrams to capture phrases or n-grams in words, the size of the features space further explodes.

In order to deal with the high dimensionality of the feature space, different routes have been explored. A first is to perform feature selection, using e.g. mutual information, and retain only features that are specific for classification (with all the potential drawbacks discussed earlier). A second approach is to use dimensionality reduction methods, such as word embeddings or LSI, and represent then documents in the much lower dimensional concept space. Finally, one may also focus on using classification algorithms that scale well with dimensionality.

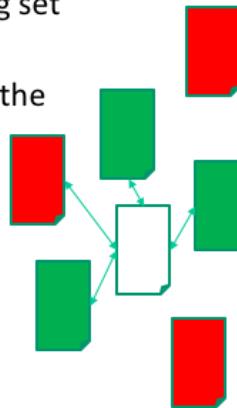
Simple Method: k-Nearest-Neighbors (kNN)

Approach: use vector space model

To classify a document D

- retrieve the k nearest neighbors in the training set according to the vector space model
- Choose the majority class label as the class of the document

If k large: $\#C / k$ estimates $P(C|D)$, the probability that the document has indeed class C



©2018, Karl Aberer, EPFL-IC, Laboratoire de systèmes d'informations répartis

Applied Classification - 72

Actually, a very straightforward method for document classification, that scales well with dimensionality, is inspired by information retrieval. In the kNN method, the top k most similar documents according to the vector space model are retrieved, and then the majority label is used as the classification of the document. Strictly speaking with this method no classifier is learnt, but the classifier consists of the vector space representation of the whole document collection, so learning is trivial in that case. The method works in practice well. One critical choice is the parameter k, that determines how large the neighborhood is that is taken into account, and can be understood as the complexity of the model. In line with what we have seen on the dependency of bias and variance on model complexity. For kNN a model is complex when k is small, since for every document very different sets of neighbors are chosen. Correspondingly it holds that for small k (complex model) we have low bias, but high variance and for large k we have high bias and low variance.

Naïve Bayes Classifier

Approach: apply Bayes law

Feature: Bag of words model: all words and their counts in the document

Using Bayes law the probability the class is C for document D

$$P(C|D) \propto P(C) \prod_{w \in D} P(w|C)$$

Another frequently used method used in document classification is Naïve Bayes Classifier. Like kNN can be understood as the analogue of classification for vector space retrieval, Naïve Bayes can be understood as the analogue of classification for probabilistic retrieval. By applying Bayes law we can derive the probability of a document belonging to a class, from two estimators (probabilities) that have to be learnt from the training data. For this method to work, words not occurring in the training set, but in the test set, need to have non-zero probabilities.

Naïve Bayes Classifier Method

Training

How characteristic is word w for class C?

$$P(w|C) = \frac{|w \in D, D \in C| + 1}{\sum_{w'} |w' \in D, D \in C| + 1}$$

How frequent is class C?

$$P(C) = \frac{|D \in C|}{|\mathcal{D}|}$$

Classification: Thus the most probable class is the one with the largest probability

$$C_{NB} = \operatorname{argmax}_C \left(\log P(C) + \sum_{w \in D} \log P(w|C) \right)$$

©2018, Karl Aberer, EPFL-IC, Laboratoire de systèmes d'informations répartis

Applied Classification - 74

For training the model we have to compute for each class how likely a word is to appear in that class and we have to compute the relative frequencies of the classes, i.e. the probability of a class label to occur. Note that in the estimation for $P(w|C)$ we add +1 to the word counts. This is called Laplacian smoothing and is necessary to avoid to assign zero probability to terms that do not occur in the training set, to avoid that subsequently a document that does contain such a term is considered as impossible when using the classifier. The reasoning is the same why smoothing has been used in probabilistic information retrieval.

Classification is performed by selecting the class with the highest probability to occur. This is computed from the logarithm of the estimated probability, as the summation is numerically more stable than multiplication.

Classification Using Word Embeddings

Reminder

government debt problems turning into banking crises as has happened in
saying that Europe needs unified banking regulation to replace the hodgepodge

↶ These words will represent *banking* ↷

Probability that word w occurs with context
word c

$$P(D = 1|w, c; \theta) = \frac{1}{1 + e^{-v_c \cdot v_w}}$$

©2018, Karl Aberer, EPFL-IC, Laboratoire de systèmes d'informations répartis

Applied Classification - 75

For document classification word embeddings can be used in different ways. One obvious way would be to use word embedding vectors as word representation, and to represent then documents as an aggregate of those vectors to obtain a low-dimensional feature space. However, another possibility is to use the word embedding approach in a more direct way, which has resulted in a classification algorithm that has been recently developed at Facebook and is known as Fasttext.

Classification Using Word Embeddings

Consider instead of (word, context) pairs,
(class, paragraph) pairs

- Word $w \rightarrow$ Class Label C
- Context words $c \rightarrow$ Paragraph p

$$\text{Learn } P(C|p) = \frac{e^{v_p \cdot v_C}}{\sum_{C'} e^{-v_p \cdot v_{C'}}} \quad (\text{SGD})$$

Paragraph feature: $v_p = \sum_{w \in p} v_w$

Then predict label from paragraph features

In this approach to document classification the idea of creating word embeddings is slightly modified. Instead of trying to predict the a word from its context, for document classification one tries to predict the class label from a text fragment, e.g. a paragraph. Technically speaking, the two problems are equivalent, as in both cases we try to predict some word from a set of words. Thus the same method for learning the word embedding can be applied, using stochastic gradient descent. Once the model is learnt, as with Naïve Bayes the class label with the highest probability to occur is chosen as prediction for the document class. If there is not enough training data, instead of learning the word vectors v_w also pre-trained vectors from other document collections can be used.

Fasttext

Classifier based on word embeddings

Extensions

- Using word n-grams (phrases)
- Subword information (character n-grams)

- Possible to build vectors for unseen words!

$$h_w = \sum_{g \in w} x_g$$

mang erai ange
man ang era gera
nge ger rai nger
Character n-grams

+ ~~ma~~ ~~erai~~

Word itself

In Fasttext different variations of the approach, using different feature sets have been explored, including the use of word phrases and character n-grams. One interesting implication of using character n-grams is that the classifier can even use feature vectors for unseen words, when they share a large number of n-grams in common with known words. The example illustrates that for example in French this can be indeed very useful.

Document Classification Summary

Widely studied problem with many applications

- Spam filtering, Sentiment Analysis, Document Search
- Increasing interest
 - Powerful methods using very large corpuses
 - Information filtering on the Internet

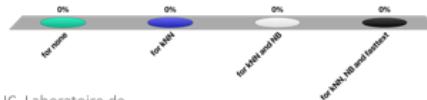
Rank	Data Mining Methods	# Papers
1	SVM	55
2	KNN	31
3	NB	23
4	ANN	10
5	Rocchio Algorithm	9
6	Association Rule Mining	4

Rank	Feature Selection Methods	# Papers
1	Chi-squared test(CHI)	21
2	Information Gain (IG)	20
3	Mutual Information(MI)	16
4	Latent Semantic Indexing (LSI), Singular Value Decomposition (SVD)	10
5	Document Frequency (DF)	8
6	Term Strength (TS)	3
7	Odds Ratio(OR)	3
8	Linear Discriminant Analysis (LDA)	3

Document classification is an area of high, both because it is required in a growing number of application domains and because at the same time the classification methods are becoming increasingly powerful due to the availability of very large document corpuses and the growing computational power available. We have discussed only a few very well known methods. The tables show the findings of a literature review on different approaches to document classification, both in terms of algorithms used and of feature selection methods applied. This gives a sense of the diversity of this field.

The dimensionality of the feature space depends on the vocabulary size ...

- A. for none
- B. for kNN
- C. for kNN and NB
- D. for kNN, NB and fasttext



©2018, Karl Aberer, EPFL-IC, Laboratoire de
systèmes d'informations répartis

References

The slides are loosely based also on:

- <http://barabasilab.neu.edu/courses/phys5116/>

Papers

- Blondel, Vincent D., et al. "Fast unfolding of communities in large networks." *Journal of statistical mechanics: theory and experiment* 2008.10 (2008): P10008
- Girvan, Michelle, and Mark EJ Newman. "Community structure in social and biological networks." *Proceedings of the national academy of sciences* 99.12 (2002): 7821-7826.

References

Document classification

- Aggarwal, Charu C., and Cheng Xiang Zhai. "A survey of text classification algorithms." *Mining text data*. Springer US, 2012. 163-222.
- Jindal, Rajni, Ruchika Malhotra, and Abha Jain. "Techniques for text classification: Literature review and current trends." *Webology* 12.2 (2015): 1
- Christopher D. Manning, Prabhakar Raghavan and Hinrich Schütze, *Introduction to Information Retrieval*, Cambridge University Press. 2008 (<http://www-nlp.stanford.edu/IR-book/>) Chapter 13
- Fasttext: <https://www.youtube.com/watch?v=CHcExDsDeHU>

Recommender Systems

- Chapter 9 in mmds.org
- Recommender Systems Handbook, Springer 2015